

29.13.implement_SGD_1.0

August 1, 2018

1 Implement SGD on bston dataset

Dataset from sklearn load_boston

The GD is $\theta = \theta - \alpha \frac{d}{d\theta}(\text{cost function})$ $equ = y = wtx + b$; $cost = (y - wtx - b)^2$ determinant wrt $w \Rightarrow -2(y - wtx - b)x$ determinant wrt $b \Rightarrow -2(y - wtx - b)$; in general : $loss$ These are in vec notation

$h = \text{np.dot}(x, \theta)$; $loss = h - y$; $cost = \text{np.sum}(loss^2) / (2m)$; $gradient = \text{np.dot}(x, loss) / m$; $\theta = \theta - \alpha \text{gradient}$

2 Objective

Implement the model and compare the intercept and weights with sklearn

3 Import data and libraries and split train and test by 80:20

```
In [139]: from sklearn.datasets import load_boston
import numpy as np
import matplotlib.pyplot as plt

boston = load_boston()
print(boston.data.shape)
boston.data[0:1]

y=np.array(boston.target)
x=np.array(boston.data)
from sklearn.preprocessing import normalize, StandardScaler
x = normalize(x, norm='l1', axis=0)
#y = normalize(y.reshape(-1, 1), norm='l1', axis=0)
#y=y.ravel()
#print("shape of y",y.shape)

x_train=x[0:450]
x_test=x[451:505]
y_train=y[0:450]
y_test=y[451:505]
```

(506, 13)

4 Create my model and cost function to calculate cost

In [140]: # Create all functions

```
#m=no of obs,n=no of features
def grad(x,y,theta,alpha=.0001,iteration=1000,k=200):

    #for SGD
    m=x.shape[0]

    idx=np.random.randint(m,size=k)
    x=x[idx,:]
    y=np.random.choice(y,k)
    #print('for SGD new size',x.shape,y.shape)

    m,n=np.shape(x)
    xt=x.transpose()
    #print("x y, theta , m,n,shape before multiply",x.shape,y.shape,theta.shape,m,n)
    oldcost=0
    bold=0
    l=0
    thetaprev=theta
    prevalpha=alpha
    optimumcost=0
    ind=0
    for i in range(0,iteration):
        l=l+1
        h=np.dot(x,theta)
        #print("\nbefore loss h and y",h.shape,y.shape)
        loss=h-y
        #print('loss',loss)
        cost=np.sum(loss**2)/(2*k)
        #if (oldcost~=0 & cost-oldcost<20):
        #    break;
        #if change is very slow increase alpha
        b=(cost-oldcost)/oldcost
        #print("At iteration %d Cost : %f oldcost : %f alpha %f cost compare %f" %(i,cost,oldcost,alpha,b))
        #if ((b>0 & bold<0) | (b<0 & bold>0)):
        if((cost>oldcost) & (oldcost>0) & (ind==0)):
            #set optimumcost only one time
            ind=1
            optimumcost=oldcost
            alpha=alpha/2
            print("opt cost,ind",optimumcost,ind)
            theta=thetaprev
            print('new alpha',alpha,'cost',cost,'old cost',oldcost)
            continue
        elif ((b>-.05) & (ind==0)):
```

```

        alpha=alpha*2
    #elif (ind==1):
    #    alpha=alpha*1.15
    #    ind=2

    gradient=np.dot(xt,loss)/k
    #print("\ngradient\n",gradient)
    theta=theta-alpha*gradient
    aa=alpha*gradient
    b=(cost-oldcost)/oldcost
    oldcost=cost
    thetaprev=theta
    prealpha=alpha
    return theta,cost

def data(x,y,k=10):
    x=np.append(x, np.ones([x.shape[0], 1]), axis=1) #with intercept
    #without intercept
    n1=np.shape(x)[1]

    theta=np.random.randn(n1)
    print("theta passed shaped",theta.shape,theta)
    newtheta,cost=grad(x,y,theta)
    print('coefficient of mymodel',newtheta)
    return newtheta,cost

def predictcost(x,newtheta):
    x=np.append(x, np.ones([x.shape[0], 1]), axis=1) #with intercept
    #without intercept
    y=np.dot(x,newtheta)
    return y

def cost1(y_pred,y_test):
    loss=-y_pred+y_test
    cost=np.sum(loss**2)/(2*y_test.shape[0])
    return cost

```

5 Use my model on boston data

```

In [141]: import warnings
          warnings.filterwarnings('ignore')

import pandas as pd
#print("train test split",x_train.shape,x_test.shape,y_train.shape,y_test.shape)
#scale=StandardScaler(with_mean=True)
#x_train=scale.fit_transform(x_train)
#x_test=scale.transform(x_test)

```

```

newtheta, cost = data(x_train, y_train)
print("my model cost from the train", cost)
#predict
y_pred = predictcost(x_train, newtheta)
cost = cost1(y_pred, y_train)


y_pred = predictcost(x_test, newtheta)
c = cost1(y_pred, y_test)

#print(y_pred.shape, y_test.shape)
print("my model cost from the test", c)
aa = pd.DataFrame({'type': ['my model'], 'train_cost': [cost], 'test_cost': [c], 'Coeff': [newtheta]})
#print(aa)

plt.scatter(y_test, y_pred)
plt.xlabel("Prices: $Y_i$")
plt.ylabel("Predicted prices: $\hat{Y}_i$")
plt.title("my model Prices vs Predicted prices: $Y_i$ vs $\hat{Y}_i$")
plt.show()

theta passed shaped (14,) [ 0.84181168  0.28353384 -1.65890218  0.67759257 -0.23250494  1.55150494
 0.95602726 -0.60950194  1.19612826  0.36614641  0.18908912  0.16858242
 1.53765322  0.13158901]
opt cost, ind 36.6582077629 1
new alpha 1.6384 cost 36.7077999525 old cost 36.6582077629
coefficient of my model [ -4.30642126e+00  4.82291220e+00 -2.55874326e+00 -1.06998034e+01
 -6.13320333e-01  1.74705492e+00 -1.84281057e-03  5.26956375e-01
 -5.46160077e-01 -2.90538283e-01  2.86636557e-01  4.13640841e-01
 9.42275633e-01  2.22545481e+01]
my model cost from the train 36.529936717
my model cost from the test 15.2824107048

```



6 Use SGDREGRESSOR model on boston data

```
In [142]: import warnings
          warnings.filterwarnings('ignore')

          # Use sklearn
          import pandas as pd
          from sklearn.linear_model import SGDRegressor
          x_train=pd.DataFrame(x_train)
          x_test=pd.DataFrame(x_test)
          y_train=pd.DataFrame(y_train)
          y_test=pd.DataFrame(y_test)
          #print(x_train)
          #print(y_train)

          print(y_train.shape,y_test.shape,x_train.shape,x_test.shape)
          lm = SGDRegressor(fit_intercept=True)
          lm.fit(x_train, y_train)
          #print(lm.coef_)
          print(lm)
          #y_test.reshape(-1,1)
          #np.reshape(y_test,54)
```

```

# print(y_test.shape)
y_pred = lm.predict(x_train)
print('coefficient of sgdregressor model', lm.intercept_, lm.coef_)

#c=cost1(y_pred,y_train)
y_predtrain = lm.predict(x_train)
print("sklearn Cost from train")
print(lm.score(x_train,y_train))

#c1=cost1(y_pred,y_test)
y_predtest = lm.predict(x_test)
print("shape", y_predtrain.shape, y_train.shape)
print(type(cost1(y_predtrain.reshape(450,1), y_train)))
print("train cost :", cost1(y_predtrain.reshape(450,1), y_train)[0])

print("test cost", cost1(y_predtest.reshape(54,1), y_test)[0])
# print(y_predtrain)
# print(y_train)
# print("sklearn Cost from test", lm.score(y_predtest, y_test))
# bb=pd.DataFrame({'type':['sgdregressor'], 'train_cost':[lm.score(x_train,y_train)], '
bb=pd.DataFrame({'type':['sgdregressor'], 'train_cost':[cost1(y_predtrain.reshape(450,1), y_train)[0]]})
aa=aa.append(bb)

plt.scatter(y_test, y_predtest)
plt.xlabel("Prices: $Y_i$")
plt.ylabel("Predicted prices: $\hat{Y}_i$")
plt.title("sgdregressor Prices vs Predicted prices: $Y_i$ vs $\hat{Y}_i$")
plt.show()


print(aa)

```

```

(450, 1) (54, 1) (450, 13) (54, 13)
SGDRegressor(alpha=0.0001, average=False, epsilon=0.1, eta0=0.01,
              fit_intercept=True, l1_ratio=0.15, learning_rate='invscaling',
              loss='squared_loss', max_iter=None, n_iter=None, penalty='l2',
              power_t=0.25, random_state=None, shuffle=True, tol=None, verbose=0,
              warm_start=False)
coefficient of sgdregressor model [ 22.73990766] [-0.03697114  0.10995223  0.01883134  0.10206
 0.03142505  0.0571567   0.01386032  0.0279774   0.03979366  0.05169953
 0.00950983]
sklearn Cost from train
-0.000943380226666
shape (450,) (450, 1)
<class 'pandas.core.series.Series'>
train cost : 45.1897757515
test cost 17.2183729505

```



	Coeff	test_cost	train_cost	\
0	[-4.30642126306, 4.82291219613, -2.55874326256...]	15.282411	45.445280	
0	[-0.0369711351004, 0.109952233902, 0.018831339...]	17.218373	45.189776	


```

type
0     mymodel
0     sgdregressor

```

7 Score comparison of mymodel and SGDRegressor

```

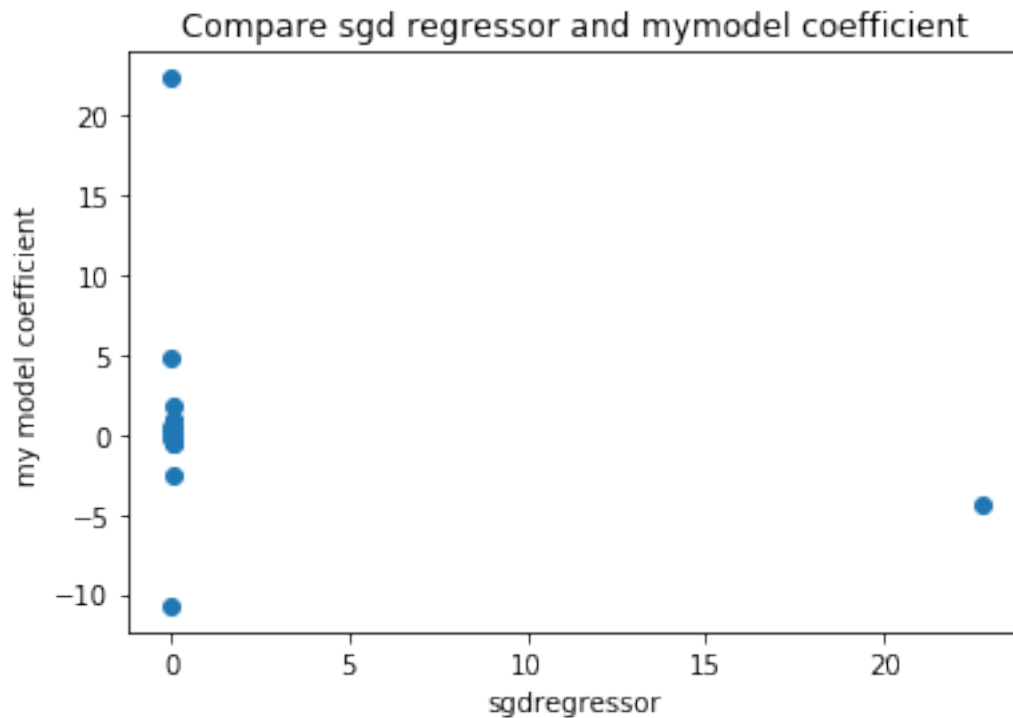
In [143]: print('sklearn coeff',lm.intercept_,lm.coef_)
          print('mymodel theta including intercept',newtheta[0],newtheta[1:])
          x1=lm.intercept_.tolist()
          x2=lm.coef_.tolist()
          x1.extend(x2)
          y=newtheta.tolist()
          import matplotlib.pyplot as plt
          plt.title("Compare sgd regressor and mymodel coefficient")
          plt.xlabel("sgdregressor")
          plt.ylabel("my model coefficient")
          plt.scatter(np.array(x1),np.array(y))
          plt.show()

```

```

sklearn coeff [ 22.73990766] [-0.03697114  0.10995223  0.01883134  0.10206679  0.03709152  0.0
0.03142505  0.0571567   0.01386032  0.0279774   0.03979366  0.05169953
0.00950983]
mymodel theta including intercept -4.30642126306 [  4.82291220e+00 -2.55874326e+00 -1.069980
1.74705492e+00 -1.84281057e-03  5.26956375e-01 -5.46160077e-01
-2.90538283e-01  2.86636557e-01  4.13640841e-01  9.42275633e-01
2.22545481e+01]

```



8 Observation

1. While trying learning rate of different value sometimes cost is going towards very high value, so tried with higher value
2. After trying more iteration cost is going down with a very small alpha value
3. sklearn model is performing better with cost close to 0 but my model cost is close to 30
4. Best cost is achieved by different trial cost function
5. theta are not changing much from initial random value (when alpha like .000005)
6. when alpha taken like .05 , gradient is becoming too high. cost function is becoming inf, theta are becoming very high
7. sklearn model is always getting global minimum but my model doesn't