

39.15.Amazon_food_review_truncated_SVD

July 14, 2018

1 Amazon food review dataset apply truncated SVD

Data set from <https://www.kaggle.com/snap/amazon-fine-food-reviews>

2 Objective

1. Take 2000 words by TFIDF importance
2. Calculate cooccurrence matrix with neighbourhood of size 5 and count how many times wi occur in context of wj
3. Then do truncated SVD
4. try multiple value of k(find optimal k by amount of variance explained)[use singular value]
5. cluster(kmeans k=50) word vector for top 2000
6. word cluster together should be related

3 Import data and libraries

```
In [97]: from sklearn.manifold import TSNE
import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer
from sklearn.cross_validation import train_test_split, KFold
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from sklearn.cross_validation import cross_val_score
from collections import Counter
```

```

from sklearn.metrics import accuracy_score
from sklearn import cross_validation
from sklearn.grid_search import GridSearchCV
from sklearn.linear_model import LogisticRegression

con = sqlite3.connect('database.sqlite')

#get only +ve and -ve review
raw_data = pd.read_sql_query("""SELECT * FROM Reviews WHERE Score != 3""", con)

```

4 Data preprocessing

```

In [98]: filtered_data=raw_data
# Score>3 a positive rating, and score<3 a negative rating.
def partition(x):
    if x < 3:
        return 'negative'
    return 'positive'

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative

filtered_data.sample(5)
filtered_data['Score'].value_counts()

#Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False)

#Deduplication of entries for same profilename,userid, time, text and take first element
sorted_data=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"})

In [206]: #take only 50000 data
print('total data \n',sorted_data['Score'].value_counts())
#clean_data=sorted_data.sample(frac=1).groupby('Score').head(10000)
#take stratified sampling i.e. positive and negative reviews are proportionate to ratio
#testing
_, clean_data = train_test_split(sorted_data, test_size = 50000, random_state=1, stratify=sorted_data['Score'])
clean_data['Score'].value_counts()

```

```

total data
positive    307063
negative    57110
Name: Score, dtype: int64

```

```

Out[206]: positive    42159

```

```
negative      7841
Name: Score, dtype: int64
```

```
In [207]: # Clean html tag and punctuation
import re
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

stop = set(stopwords.words('english')) #set of stopwords
sno = nltk.stem.SnowballStemmer('english') #initialising the snowball stemmer

#substitute html tag and punctuation
def cleanhtml(sentence): #function to clean the word of any html-tags
    cleanr = re.compile('<.*?>')
    cleantext = re.sub(cleanr, ' ', sentence)
    return cleantext
def cleanpunc(sentence): #function to clean the word of any punctuation or special c
    cleaned = re.sub(r'[?|!|\'|\"|#]',r'',sentence)
    cleaned = re.sub(r'[.,|)|(|\\|/]',r' ',cleaned)
    return cleaned
print(sno.stem('tasty'))

i=0
str1=' '
mystop={'of','four','one','would'}
final_string=[]
all_positive_words=[] # store words from +ve reviews here
all_negative_words=[] # store words from -ve reviews here.
s=''
#Create new catagory as Cleanedtext after removing htmltag and punctuation and upper
for sent in clean_data['Text'].values:
    filtered_sentence=[]
    sent=cleanhtml(sent) # remove HTML tags
    for w in sent.split():
        for cleaned_words in cleanpunc(w).split():
            if((cleaned_words.isalpha()) & (len(cleaned_words)>2)):
                if((cleaned_words.lower() not in stop) & (cleaned_words.lower() not
                    s=(sno.stem(cleaned_words.lower())).encode('utf8')
                    filtered_sentence.append(s)
                    if (clean_data['Score'].values)[i] == 'positive':
                        all_positive_words.append(s) #list of all words used to desc
                    if(clean_data['Score'].values)[i] == 'negative':
                        all_negative_words.append(s) #list of all words used to desc
            else:
                continue
        else:
            continue
```

```

        continue
    str1 = b" ".join(filtered_sentence) #final string of cleaned words

    final_string.append(str1)
    i+=1

clean_data['CleanedText']=final_string
print(clean_data.shape)
#Sort data on timestamp
clean_data=clean_data.sort_values(by=['Time'],ascending=False)
#clean_data
clean_data['CleanedText'].sample(2)
clean_data['CleanedText'].iloc[0]

tasti
(50000, 11)

```

C:\Users\suman\Anaconda3\lib\site-packages\ipykernel_launcher.py:52: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

Out[207]: b'use lot coconut late granola cake cup cake etc particular brand serv purpos well p

5 Get top 2000 words by TFIDF score and create co-occurrence matrix by window 5

```

In [209]: x=clean_data['CleanedText'].values
          y = clean_data['Score']
          tf_idf_vect = TfidfVectorizer()
          final_counts = tf_idf_vect.fit_transform(x)
          #use the same vectors to convert test data
          indices = np.argsort(tf_idf_vect.idf_)[::-1]
          features = tf_idf_vect.get_feature_names()
          #testing
          top_n = 2000
          top_features = [features[i] for i in indices[:top_n]]
          print (top_features[0:20])
          print('len of top feature',len(top_features))

          #remove other words from review
          final_string=[]
          all_string=[]

          i=0

```

```

for sent in clean_data['CleanedText'].values:

    filtered_sentence=[]
    for w in sent.decode('utf8').split():
        if(w in top_features):
            filtered_sentence.append(w.encode('utf8'))

        else:
            continue
    str1 = b" ".join(filtered_sentence) #final string of cleaned words
    if ((i<5) & (str1!=b'')):
        print('sentence copy',str1)
    if (str1.decode('utf8') !=''):
        i=i+1
        final_string.append(str1)
        all_string.append(str1)

#clean_data['CleanedText']=final_string
#Now final_string is ready to work with
#print(clean_data['CleanedText'].shape)
all_string[0:20]

```

```

['çaykur', 'jail', 'jaegermeister', 'jagar', 'jager', 'jagermeister', 'jagger', 'jaguar', 'jaim',
len of top feature 2000
sentence copy b'howevert'
sentence copy b'jeesh'
sentence copy b'kadoda'
sentence copy b'hula'
sentence copy b'gough galantin'

```

```

Out[209]: [b'howevert',
b'jeesh',
b'kadoda',
b'hula',
b'gough galantin',
b'gami',
b'likelt',
b'gallopin',
b'krapelien holm',
b'lakritz',
b'gayelord',
b'fuggedaboud',
b'jeweltim jeweltim',
b'happitud geen',
b'hummingbyrd',
b'gnash',
b'jagermeister',

```

```

b'gravey',
b'intersess keiki',
b'ingrededi']

```

```

In [210]: #Convert to cooccurrence mat
          #type(final_string)
          print(len(top_features))
          window=5
          len1=len(top_features)
          #print(len)
          m=np.zeros([len1,len1])
          columns=top_features
          rows=top_features
          df=pd.DataFrame(m,columns=columns,index=rows)
          #print(df)

          def cal_occ(sentence,df):
              sen=sentence.split()
              l=len(sen)-1
              for i,word in enumerate(sen):
                  #loop through every sentence in a window and get neighest words and keep a
                  for j in range(max(i-window,0),min(i+window+1,l+1)):
                      if word!=sen[j]:
                          #print('printing',word,sen[j])
                          df[word][sen[j]]+=1

              for sentence in final_string:
                  #print('call',sentence)
                  cal_occ(sentence.decode('utf8'),df)

          print(df.shape)

2000
(2000, 2000)

```

```

In [211]: from sklearn.preprocessing import StandardScaler
          #Get cooccurring words for a given word
          print('The most similar word like ',df.index[4])
          aa=df.iloc[4]
          bb=aa.sort_values(ascending=False)
          print(type(bb))

```

```

The most similar word like  jager
<class 'pandas.core.series.Series'>

```

```

In [212]: #df

```

6 Create countvectorizer using cooccurence matrix

```
In [213]: #count_vect = CountVectorizer(vocabulary=top_features) #in scikit-learn
          #X = count_vect.fit_transform(final_string)

          #print(X.shape)

          #Cooccurence matrix
          #X = (X.T * X) # this is co-occurrence matrix in sparse csr format
          #X.setdiag(0) # sometimes you want to fill same word cooccurence to 0
          #print(X.todense())

          #print(count_vect.vocabulary_)

          #Create truncated SVD
          from sklearn.decomposition import TruncatedSVD

          #Try different component
          l=[20,50,100,150,200,250]
          for i in l:
              svd = TruncatedSVD(n_components=i, n_iter=7, random_state=0)
              svd.fit(df.values)
              #print(svd.explained_variance_ratio_)
              l1=svd.explained_variance_ratio_
              print('% variance explained with component ',i,svd.explained_variance_ratio_.sum())
              #print('singular values',svd.singular_values_)

          #So looks like with 25 component 96% variance is explained

% variance explained with component  20 0.614688953311
% variance explained with component  50 0.775927938321
% variance explained with component 100 0.862135844503
% variance explained with component 150 0.889968288638
% variance explained with component 200 0.91780506569
% variance explained with component 250 0.945634269034
```

SO by 250 component 95% variance is explained # Use SVD

```
In [229]: #VT = svd.components_
          #TruncatedSVD is basically a wrapper around sklearn.utils.extmath.randomized_svd; yo

          from sklearn.utils.extmath import randomized_svd

          U, Sigma, VT = randomized_svd(df.values,
                                       n_components=250,
                                       n_iter=50,
                                       random_state=0)
```

```

print('U value\n')
#print(U)
print('sigma value\n')
#print(Sigma)
print('VT value\n')
#print(VT)
print(U.shape,Sigma.shape,VT.shape)
print('1st word vector representation',df.index[0])

```

U value

sigma value

VT value

(2000, 250) (250,) (250, 2000)
 1st word vector representation çaykur

7 Form cluster of 10 using those important words SVD value

8 Get top few words similar to a random word

```

In [230]: from random import randint
          j=randint(0, 2000)
          print('1st word vector representation',df.index[j],' for j',j)#print(U.shape[0])
          #Calculate distance of this word with all words and sort in descending order
          #take log transform
          U=np.log(U+1)
          l=[]
          for i in range(U.shape[1]):
              a=np.linalg.norm(U[j]-U[i])
              l.append(a)
          l1=sorted(range(len(l)), key=lambda k: l[k])
          #print(l)
          print('top 10 words similar to ',df.index[j],' are ',df.index[l1[0:10]])

```

1st word vector representation granmder for j 1331
 top 10 words similar to granmder are Index(['jitterbean', 'jot', 'josephus', 'jedi', 'jeera',
 'jitterbug', 'jem', 'jolen'],
 dtype='object')

```

In [233]: from sklearn.cluster import KMeans
          # Now U is vec presentation of words
          n_clusters=10
          kmeans=KMeans(n_clusters=10, random_state=0).fit(U)
          kmeans.cluster_centers_

```




```

sse = {}
for k in range(1, 10):
    kmeans = KMeans(init='k-means++', n_clusters=k, max_iter=100).fit(U)
    sse[k] = kmeans.inertia_ # Inertia: Sum of distances of samples to their closest
plt.figure()
plt.plot(list(sse.keys()), list(sse.values()))
plt.title("Plot of loss against k value using BOW")
plt.xlabel("Number of cluster")
plt.ylabel("SSE")
plt.show()

#a=np.where(kmeans.labels_ == 1)[0]
#b=np.where(kmeans.labels_ == 0)[0]
#check 5 text for cluster 1

kmeans = KMeans(init='k-means++', n_clusters=50, max_iter=100).fit(U)
n_clusters=50
print(a.shape)
for i in range(n_clusters):
    a=np.where(kmeans.labels_ == i)[0]
    print('in cluster \n',i)
    print(a[0:10])
    k=0
    for j in a:
        k=k+1
        if (k<10):
            print(top_features[j][:10])

```



```
(1,)
in cluster
0
[0 1 2 3 4 5 6 7 8 9]
çaykur
jail
jaegermeis
jagar
jager
jagermeist
jagger
jaguar
jaim
in cluster
1
[882]
kittredg
in cluster
2
[1289]
gradul
in cluster
3
[968]
```

kelsey
in cluster
4
[22]
jab
in cluster
5
[1845]
hushpuppi
in cluster
6
[1892]
hotown
in cluster
7
[632]
lanolin
in cluster
8
[1570]
gingerspice
in cluster
9
[1966]
hairstyl
in cluster
10
[1063]
knowlton
in cluster
11
[1213]
gleam
in cluster
12
[1280]
grenad
in cluster
13
[1573]
giovanni
in cluster
14
[984]
kwazulu
in cluster
15
[678]

laurenc
in cluster
16
[123]
istelf
in cluster
17
[1017 1806]
laevulos
higer
in cluster
18
[1661]
genui
in cluster
19
[1654]
genovo
in cluster
20
[1636]
genom
in cluster
21
[1733]
homag
in cluster
22
[987]
kushka
in cluster
23
[827]
lindo
in cluster
24
[654 1665]
lamma
genteel
in cluster
25
[1863]
hydron
in cluster
26
[1440]
gaul
in cluster

27
[992]
kuechenmei
in cluster
28
[281]
implant
in cluster
29
[776]
lisey
in cluster
30
[1589]
gim
in cluster
31
[1385]
gravita
in cluster
32
[1078]
komissbrot
in cluster
33
[160]
juarez
in cluster
34
[1692]
hobbits
in cluster
35
[1341]
grandaroma
in cluster
36
[387]
immacul
in cluster
37
[373]
iceberg
in cluster
38
[1487]
fue
in cluster

39
[1400]
ganoderma
in cluster
40
[1797]
hijiki
in cluster
41
[448]
insalata
in cluster
42
[798]
licken
in cluster
43
[1796]
hijo
in cluster
44
[101]
ipodo
in cluster
45
[86]
ironwork
in cluster
46
[1430]
gastroente
in cluster
47
[685 963]
latitud
kensington
in cluster
48
[537]
ingredien
in cluster
49
[841]
kiefer

9 Observation

Most of the cluster contains 1-2 words and most of the words in one cluster To find the similar word its not giving proper words, maybe non english words and text cleanning is required much

In []:

Ignore the above 2 plots those plots are plotted below again