

32.19.Amazon_food_review_GBDT

July 1, 2018

1 Amazon food review dataset apply GBDT to predict polarity review

Data set from <https://www.kaggle.com/snap/amazon-fine-food-reviews>

2 Objective

Here try BOW,TFIDF,avg W2V, avg TFIDFW2V and different hyperparameter GBDT(no of base learner,depth,learning rate)

Import data and libraries

```
In [15]: import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer
con = sqlite3.connect('database.sqlite')
#get only +ve and -ve review
raw_data = pd.read_sql_query("""SELECT * FROM Reviews WHERE Score != 3""", con)
```

3 Data preprocessing

```
In [16]: filtered_data=raw_data
# Score>3 a positive rating, and score<3 a negative rating.
def partition(x):
    if x < 3:
        return 'negative'
    return 'positive'
```

```

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
#filtered_data.sample(5)
filtered_data['Score'].value_counts()
#Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False)
#Deduplication of entries for same profilename,userid, time, text and take first element
sorted_data=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"})

```

```

In [17]: #take only 2500 + 2500 data
from sklearn.cross_validation import train_test_split,KFold
_, clean_data = train_test_split(sorted_data, test_size = 10000, random_state=0, stratify=sorted_data['Score'])
clean_data['Score'].value_counts()

```

```

Out[17]: positive      8432
         negative      1568
         Name: Score, dtype: int64

```

```

In [18]: # Clean html tag and punctuation
import re
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

stop = set(stopwords.words('english')) #set of stopwords
sno = nltk.stem.SnowballStemmer('english') #initialising the snowball stemmer
#substitute html tag and punctuation
def cleanhtml(sentence): #function to clean the word of any html-tags
    cleanr = re.compile('<.*?>')
    cleantext = re.sub(cleanr, ' ', sentence)
    return cleantext
def cleanpunc(sentence): #function to clean the word of any punctuation or special character
    cleaned = re.sub(r'[?|!|\'|\"|#]',r'',sentence)
    cleaned = re.sub(r'[.,|)|(|\\|/]',r'',cleaned)
    return cleaned
#print(sno.stem('tasty'))

i=0
str1=' '
final_string=[]
all_positive_words=[] # store words from +ve reviews here
all_negative_words=[] # store words from -ve reviews here.
s=''
#Create new category as Cleanedtext after removing htmltag and punctuation and uppercasing

```

```

for sent in clean_data['Text'].values:
    filtered_sentence=[]
    #print(sent);
    sent=cleanhtml(sent) # remove HTML tags
    for w in sent.split():
        for cleaned_words in cleanpunc(w).split():
            if((cleaned_words.isalpha()) & (len(cleaned_words)>2)):
                if(cleaned_words.lower() not in stop):
                    s=(sno.stem(cleaned_words.lower())).encode('utf8')
                    filtered_sentence.append(s)
                    if (clean_data['Score'].values)[i] == 'positive':
                        all_positive_words.append(s) #list of all words used to descr
                    if(clean_data['Score'].values)[i] == 'negative':
                        all_negative_words.append(s) #list of all words used to descr
                else:
                    continue
            else:
                continue
    str1 = b" ".join(filtered_sentence) #final string of cleaned words

    final_string.append(str1)
    i+=1

```

```

clean_data['CleanedText']=final_string
clean_data.shape
#Sort data on timestamp
clean_data=clean_data.sort_values(by=['Time'],ascending=False)
clean_data.sample(2)

```

```

Out[18]:
      Id  ProductId  UserId  ProfileName \
69317  75358  B002YR97QA  AJ1CD70I9YFFZ  Mary Frost "mrtshome"
163375  177139  B005ZAULI6  A2TG5JHF4CCPGY  Jacqueline A. Johnson

      HelpfulnessNumerator  HelpfulnessDenominator  Score  Time \
69317                    0                      0  positive  1327795200
163375                    0                      0  positive  1328918400

      Summary \
69317  Teatulia tea with neem
163375  We will see.

      Text \
69317  This tea is interesting. My friend says neem ...
163375  I have not had any of this tea yet as it only ...

      CleanedText
69317  b'tea interest friend say neem tast take get u...
163375  b'tea yet arriv today told compani would recei...

```

4 Split train and test data

```
In [19]: x=clean_data['CleanedText'].values
        y = clean_data['Score']
        n=x.shape[0]
        n1=int(n*.3)
        X_test_raw = x[0:n1]
        X_train_raw= x[n1:n+1]
        y_test=y[0:n1]
        y_train=y[n1:n+1]
        print('size of X_train, X_test, y_train , y_test ',X_train_raw.shape, X_test_raw.shape, y_train.shape, y_test.shape)
        print("positive and negative review in train and test\n",y_train.value_counts(),"\n",y_test.value_counts())

size of X_train, X_test, y_train , y_test (7000,) (3000,) (7000,) (3000,)
positive and negative review in train and test
positive    5939
negative    1061
Name: Score, dtype: int64
positive    2493
negative     507
Name: Score, dtype: int64
```

5 Use GDDT

6 BOW

```
In [20]: #now convert CleanedText to TDM
        count_vect = CountVectorizer() #in scikit-learn
        X_train = count_vect.fit_transform(X_train_raw)

        #use the same vectors to convert test data
        X_test=count_vect.transform(X_test_raw)
        print(X_train.get_shape(),X_test.get_shape())
        X_train_BOW=X_train
        X_test_BOW=X_test

        #print(final_counts[0,:])# this is stored like dict format only non zero values. sparse matrix
        #x = pd.DataFrame(final_counts.toarray())#this is stored like dataframe format all 0 values
        # sparse matrix in csr format works faster compare to dense format
        #print(x.shape,x.loc[0])# sparse matrix in csr format works faster compare to dense format
        #print(x.shape,x.loc[0])

(7000, 11134) (3000, 11134)
```

```
In [21]: from sklearn.cross_validation import train_test_split
        from sklearn.neighbors import KNeighborsClassifier
```

```

from sklearn.metrics import accuracy_score
from sklearn.cross_validation import cross_val_score
from collections import Counter
from sklearn.metrics import accuracy_score
from sklearn import cross_validation

# Use GBDT
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.grid_search import GridSearchCV
learning_rate=[.9,.8,.1,.2]
n_estimators=[10,20,30]
max_depth=[3,4,5,6,7,8]
tuned_parameters=dict(learning_rate=learning_rate,n_estimators=n_estimators,max_depth=
#Using GridSearchCV
model = GridSearchCV(GradientBoostingClassifier(), tuned_parameters, cv=5)
model.fit(X_train, y_train)
print('Best parameters \n',model.best_estimator_)
optimumlearning_rate=model.best_estimator_.learning_rate
optimumn_estimators=model.best_estimator_.n_estimators
optimummax_depth=model.best_estimator_.max_depth
#build model with optimum parameter
model = GradientBoostingClassifier(learning_rate=optimumlearning_rate,n_estimators=optimumn_estimators,max_depth=optimummax_depth)
model.fit(X_train, y_train)

#Store scores
pred=model.predict(X_test)
mat=pd.crosstab(y_test, pred, rownames=['Actual'], colnames=['Predicted'], margins=True)
tp=mat.iloc[1,1] ;tn=mat.iloc[0,0] ;fp=mat.iloc[0,1];fn=mat.iloc[1,0];precision=tp/(tp+fn);recall=tp/(tp+fn);fscoretest=2*precision*recall/(precision+recall)

pred=model.predict(X_train)
mat=pd.crosstab(y_train, pred, rownames=['Actual'], colnames=['Predicted'], margins=True)
print(mat);tp=mat.iloc[1,1] ;tn=mat.iloc[0,0];fp=mat.iloc[0,1] ;fn=mat.iloc[1,0] ;precision=tp/(tp+fn);recall=tp/(tp+fn);fscoretrain=2*precision*recall/(precision+recall)

aa=pd.DataFrame({'type':['BOW GDDT'],'depth':[optimummax_depth],'estimator':[optimumn_estimators],'learningrate':[optimumlearning_rate],'accuracy_train':[model.score(X_train,y_train)],'fscore_train':[fscoretrain],'accuracy_test':[model.score(X_test,y_test)]})

```

Best parameters

```

GradientBoostingClassifier(criterion='friedman_mse', init=None,
                           learning_rate=0.2, loss='deviance', max_depth=7,
                           max_features=None, max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,

```

```

        min_weight_fraction_leaf=0.0, n_estimators=30,
        presort='auto', random_state=None, subsample=1.0, verbose=0,
        warm_start=False)
Predicted  negative  positive  All
Actual
negative      601      460  1061
positive       4     5935  5939
All           605     6395  7000

```

7 TFIDF

```

In [22]: tf_idf_vect = TfidfVectorizer()
         final_counts = tf_idf_vect.fit_transform(X_train_raw)
         #use the same vectors to convert test data
         X_test=count_vect.transform(X_test_raw)
         X_train_tfidf=X_train
         X_test_tfidf=X_test

         print(X_train.get_shape(),X_test.get_shape())
         #Use scale of train and apply to test
         from sklearn.preprocessing import StandardScaler
         scaler = StandardScaler(with_mean=False).fit(X_train)
         X_train = scaler.transform(X_train)
         X_test  = scaler.transform(X_test)
         print('size of X_train, X_test, y_train , y_test ',X_train.shape, X_test.shape,y_train

         learning_rate=[.9,.8,.1,.2]
         n_estimators=[10,20,30]
         max_depth=[3,4,5,6,7,8]
         tuned_parameters=dict(learning_rate=learning_rate,n_estimators=n_estimators,max_depth=
         #Using GridSearchCV
         model = GridSearchCV(GradientBoostingClassifier(), tuned_parameters, cv=5)
         model.fit(X_train, y_train)
         print('Best parameters \n',model.best_estimator_)
         optimumlearning_rate=model.best_estimator_.learning_rate
         optimumn_estimators=model.best_estimator_.n_estimators
         optimummax_depth=model.best_estimator_.max_depth
         #build model with optimum parameter
         model = GradientBoostingClassifier(learning_rate=optimumlearning_rate,n_estimators=op
         max_depth=optimummax_depth)
         model.fit(X_train, y_train)

         #Store scores
         pred=model.predict(X_test)
         mat=pd.crosstab(y_test, pred, rownames=['Actual'], colnames=['Predicted'], margins=Tr

```

```

tp=mat.iloc[1,1] ;tn=mat.iloc[0,0] ;fp=mat.iloc[0,1];fn=mat.iloc[1,0];precision=tp/(tp+fp)
recall=tp/(tp+fn);fscoretest=2*precision*recall/(precision+recall)

pred=model.predict(X_train)
mat=pd.crosstab(y_train, pred, rownames=['Actual'], colnames=['Predicted'], margins=True)
print(mat);tp=mat.iloc[1,1] ;tn=mat.iloc[0,0];fp=mat.iloc[0,1] ;fn=mat.iloc[1,0] ;precision=tp/(tp+fp)
recall=tp/(tp+fn);fscoretrain=2*precision*recall/(precision+recall)

bb=pd.DataFrame({'type':['TFIDF GDDT'],'depth':[optimummax_depth],'estimator':[optimumestimator],
                 'learningrate':[optimumlearning_rate],'accuracy_train':[model.score(X_train,y_train)],
                 'fscore_train':[fscoretrain],'accuracy_test':[model.score(X_test,y_test)]})
aa=aa.append(bb)

```

(7000, 11134) (3000, 11134)
size of X_train, X_test, y_train , y_test (7000, 11134) (3000, 11134) (7000,) (3000,)

Best parameters

```

GradientBoostingClassifier(criterion='friedman_mse', init=None,
                           learning_rate=0.2, loss='deviance', max_depth=7,
                           max_features=None, max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=30,
                           presort='auto', random_state=None, subsample=1.0, verbose=0,
                           warm_start=False)

```

Predicted \ Actual	negative	positive	All
negative	602	459	1061
positive	5	5934	5939
All	607	6393	7000

8 AVG W2V

```

In [23]: #ignore warning
import warnings
warnings.filterwarnings('ignore')
from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle
model = KeyedVectors.load_word2vec_format('GoogleNews-vectors-negative300.bin.gz', binary=True)
import gensim
i=0
#create a list of list to be used in W2V
list_of_sent_train=[]
for sent in X_train_raw: #clean_data['CleanedText'].values:
    filtered_sentence=[]

```

```

    #sent=cleanhtml(sent)
    for w in sent.split():
        #for cleaned_words in cleanpunc(w).split():
        for cleaned_words in w.split():
            if(cleaned_words.isalpha()):
                filtered_sentence.append(cleaned_words.lower().decode('utf8'))
            else:
                continue
    list_of_sent_train.append(filtered_sentence)
#convert each sentence's words to a vector of 50 dimension. Dont construct vec if wor
#and 4 core processor
w2v_model=gensim.models.Word2Vec(list_of_sent_train,min_count=5,size=50, workers=4)

# average Word2Vec
# for each sentence make average of vectors by (vectors of each words)/(total no of w
# compute average word2vec for each review.
sent_vectors_train = []; # the avg-w2v for each sentence/review is stored in this lis
for sent in list_of_sent_train: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    cnt_words =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        try:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
        except:
            pass
    sent_vec /= cnt_words
    sent_vectors_train.append(sent_vec)

#convert W2V test data
i=0
#create a list of list to be used in W2V
list_of_sent_test=[]
for sent in X_test_raw: #clean_data['CleanedText'].values:
    filtered_sentence=[]
    #sent=cleanhtml(sent)
    for w in sent.split():
        #for cleaned_words in cleanpunc(w).split():
        for cleaned_words in w.split():
            if(cleaned_words.isalpha()):
                filtered_sentence.append(cleaned_words.lower().decode('utf8'))
            else:
                continue
    list_of_sent_test.append(filtered_sentence)
#convert each sentence's words to a vector of 50 dimension. Dont construct vec if wor
#and 4 core processor

```



```

w2v_model=gensim.models.Word2Vec(list_of_sent_test,min_count=5,size=50, workers=4)

# average Word2Vec
# for each sentence make average of vectors by (vectors of each words)/(total no of words)
# compute average word2vec for each review.
sent_vectors_test = []; # the avg-w2v for each sentence/review is stored in this list
for sent in list_of_sent_test: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    cnt_words =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        try:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
        except:
            pass
    sent_vec /= cnt_words
    sent_vectors_test.append(sent_vec)

# try
X_train = pd.DataFrame(sent_vectors_train)
X_test = pd.DataFrame(sent_vectors_test)
X_train_w2v=X_train
X_test_w2v=X_test
print('size of X_train, X_test, y_train , y_test ',X_train.shape, X_test.shape,y_train.shape,y_test.shape)

# Use GBDT
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.grid_search import GridSearchCV

#GradientBoostingClassifier(learning_rate=0.1, n_estimators=100, max_depth=3, random_state=0)

learning_rate=[.9,.8,.1,.2]
n_estimators=[8,12,15]
max_depth=[3,4,5,6,7,8]

tuned_parameters=dict(learning_rate=learning_rate,n_estimators=n_estimators,max_depth=max_depth)

#Using GridSearchCV
model = GridSearchCV(GradientBoostingClassifier(), tuned_parameters, cv=5)
model.fit(X_train, y_train)

print('Best parameters \n',model.best_estimator_)
#print('Model test score', model.score(X_test, y_test))

optimumlearning_rate=model.best_estimator_.learning_rate
optimumn_estimators=model.best_estimator_.n_estimators

```

```

optimummax_depth=model.best_estimator_.max_depth

#build model with optimum parameter
model = GradientBoostingClassifier(learning_rate=optimumlearning_rate,n_estimators=optimumn_estimators,
                                   max_depth=optimummax_depth)
model.fit(X_train, y_train)

#Store scores
pred=model.predict(X_test)
mat=pd.crosstab(y_test, pred, rownames=['Actual'], colnames=['Predicted'], margins=True)
tp=mat.iloc[1,1] ;tn=mat.iloc[0,0] ;fp=mat.iloc[0,1];fn=mat.iloc[1,0];precision=tp/(tp+fp)
recall=tp/(tp+fn);fscoretest=2*precision*recall/(precision+recall)

pred=model.predict(X_train)
mat=pd.crosstab(y_train, pred, rownames=['Actual'], colnames=['Predicted'], margins=True)
print(mat);tp=mat.iloc[1,1] ;tn=mat.iloc[0,0];fp=mat.iloc[0,1] ;fn=mat.iloc[1,0] ;precision=tp/(tp+fp)
recall=tp/(tp+fn);fscoretrain=2*precision*recall/(precision+recall)

bb=pd.DataFrame({'type':['AVG W2V GDDT'],'depth':[optimummax_depth],'estimator':[optimumestimator],
                 'learningrate':[optimumlearning_rate],'accuracy_train':[model.score(X_train,y_train)],
                 'fscore_train':[fscoretrain],'accuracy_test':[model.score(X_test,y_test)],
                 'fscore_test':[fscoretest]})
aa=aa.append(bb)

```

size of X_train, X_test, y_train , y_test (7000, 50) (3000, 50) (7000,) (3000,)
 Best parameters

```

GradientBoostingClassifier(criterion='friedman_mse', init=None,
                           learning_rate=0.2, loss='deviance', max_depth=4,
                           max_features=None, max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=12,
                           presort='auto', random_state=None, subsample=1.0, verbose=0,
                           warm_start=False)

```

Predicted \ Actual	negative	positive	All
negative	57	1004	1061
positive	3	5936	5939
All	60	6940	7000

9 AVG TFIDF W2V

```

In [24]: #ignore warning
import warnings
warnings.filterwarnings('ignore')

from gensim.models import Word2Vec

```

```

from gensim.models import KeyedVectors
import pickle
model = KeyedVectors.load_word2vec_format('GoogleNews-vectors-negative300.bin.gz', binary=True)
import gensim

tf_idf_vect = TfidfVectorizer()
final_tf_idf=tf_idf_vect.fit_transform(X_train_raw)
tfidf_feat = tf_idf_vect.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectors_train = []; # the tfidf-w2v for each sentence/review is stored in row=0;
#calculate avg tfidf score for each sentences
for sent in list_of_sent_train: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        try:
            vec = w2v_model.wv[word]#calculate w2v for each word
            # obtain the tf_idfidf of a word in a sentence/review
            tf_idf = final_tf_idf[row, tfidf_feat.index(word)]#get tfidf score of each word
            sent_vec += (vec * tf_idf) # multiply vec with tfidf of each word and cumulate
            weight_sum += tf_idf # also add tfidf sums in each sentence
        except:
            pass
    sent_vec /= weight_sum
    tfidf_sent_vectors_train.append(sent_vec)
    row += 1
#tfidf_sent_vectors.

# do for test
final_tf_idf=tf_idf_vect.transform(X_test_raw)
tfidf_sent_vectors_test = []; # the tfidf-w2v for each sentence/review is stored in row=0;
#calculate avg tfidf score for each sentences
for sent in list_of_sent_test: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        try:
            vec = w2v_model.wv[word]#calculate w2v for each word
            # obtain the tf_idfidf of a word in a sentence/review
            tf_idf = final_tf_idf[row, tfidf_feat.index(word)]#get tfidf score of each word
            sent_vec += (vec * tf_idf) # multiply vec with tfidf of each word and cumulate
            weight_sum += tf_idf # also add tfidf sums in each sentence
        except:
            pass
    sent_vec /= weight_sum

```

```

        tfidf_sent_vectors_test.append(sent_vec)
        row += 1

X_train = pd.DataFrame(tfidf_sent_vectors_train)
X_test = pd.DataFrame(tfidf_sent_vectors_test)
X_train_w2vtfidf=X_train
X_test_w2vtfidf=X_test

print('size of X_train, X_test, y_train , y_test ',X_train.shape, X_test.shape,y_train.shape)

# Use GBDT
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.grid_search import GridSearchCV

#GradientBoostingClassifier(learning_rate=0.1, n_estimators=100, max_depth=3, random_state=0)

learning_rate=[.9,.8,.2]
n_estimators=[5,8,10]
max_depth=[3,4,5,6,7]

tuned_parameters=dict(learning_rate=learning_rate,n_estimators=n_estimators,max_depth=max_depth)

#Using GridSearchCV
model = GridSearchCV(GradientBoostingClassifier(), tuned_parameters, cv=5)
model.fit(X_train, y_train)

print('Best parameters \n',model.best_estimator_)
#print('Model test score', model.score(X_test, y_test))

optimumlearning_rate=model.best_estimator_.learning_rate
optimumn_estimators=model.best_estimator_.n_estimators
optimummax_depth=model.best_estimator_.max_depth

#build model with optimum parameter
model = GradientBoostingClassifier(learning_rate=optimumlearning_rate,n_estimators=optimumn_estimators,
                                   max_depth=optimummax_depth,random_state=0)
model.fit(X_train, y_train)

#Store scores
pred=model.predict(X_test)
mat=pd.crosstab(y_test, pred, rownames=['Actual'], colnames=['Predicted'], margins=True)
tp=mat.iloc[1,1] ;tn=mat.iloc[0,0] ;fp=mat.iloc[0,1];fn=mat.iloc[1,0];precision=tp/(tp+fn)
recall=tp/(tp+fn);fscoretest=2*precision*recall/(precision+recall)

pred=model.predict(X_train)

```

```
mat=pd.crosstab(y_train, pred, rownames=['Actual'], colnames=['Predicted'], margins=True)
print(mat);tp=mat.iloc[1,1] ;tn=mat.iloc[0,0];fp=mat.iloc[0,1] ;fn=mat.iloc[1,0] ;pre
recall=tp/(tp+fn);fscoretrain=2*precision*recall/(precision+recall)
```

```
bb=pd.DataFrame({'type':['AVG W2V TFIDF GBDT'],'depth':[optimummax_depth],'estimator'
                 'learningrate':[optimumlearning_rate],'accuracy_train':[model.score
                 'fscore_train':[fscoretrain],'accuracy_test':[model.score(X_test,y_t
aa=aa.append(bb)
```

size of X_train, X_test, y_train , y_test (7000, 50) (3000, 50) (7000,) (3000,)
Best parameters

```
GradientBoostingClassifier(criterion='friedman_mse', init=None,
                           learning_rate=0.2, loss='deviance', max_depth=3,
                           max_features=None, max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=5, presort='auto',
                           random_state=None, subsample=1.0, verbose=0,
                           warm_start=False)
```

Predicted \ Actual	negative	positive	All
negative	2	1059	1061
positive	0	5939	5939
All	2	6998	7000

10 plot for different values and accuracy

```
In [25]: import warnings
warnings.filterwarnings('ignore')
from sklearn.model_selection import validation_curve
#create plot for training and test validation
# We cannot put 0 in parameter it will give error while validation_curve
learning_rate=[.9,.8,.1,.2]
n_estimators=[5,8,10]
max_depth=[3,4,5,6,7]
param_range=[.9,.8,.1,.2]


train_scores, test_scores = validation_curve(GradientBoostingClassifier(), X_train, y,
                                             param_range=learning_rate,cv=5)

#print(train_scores,test_scores)
train_scores_mean = np.mean(train_scores, axis=1)
train_scores_std = np.std(train_scores, axis=1)
test_scores_mean = np.mean(test_scores, axis=1)
test_scores_std = np.std(test_scores, axis=1)
```

```

plt.plot(param_range, train_scores_mean, label="Training score", color="black")
plt.plot(param_range, test_scores_mean, label="Cross-validation score", color="red")
plt.fill_between(param_range, train_scores_mean - train_scores_std, train_scores_mean + train_scores_std, color="black")
plt.fill_between(param_range, test_scores_mean - test_scores_std, test_scores_mean + test_scores_std, color="red")
plt.title("Validation Curve for GBDT for different learning rate")
plt.xlabel("learning rate")
plt.ylabel("Accuracy Score")
plt.xlim(.1,.9)
plt.tight_layout()
plt.legend(loc="best")
plt.show()

```



```

In [26]: import warnings
warnings.filterwarnings('ignore')
from sklearn.model_selection import validation_curve
#create plot for training and test validation
# We cannot put 0 in parameter it will give error while validation_curve
learning_rate=[.9,.8,.1,.2]
n_estimators=[5,8,10]
max_depth=[3,4,5,6,7]
param_range=[5,8,10]

train_scores, test_scores = validation_curve(GradientBoostingClassifier(), X_train, y_train,

```

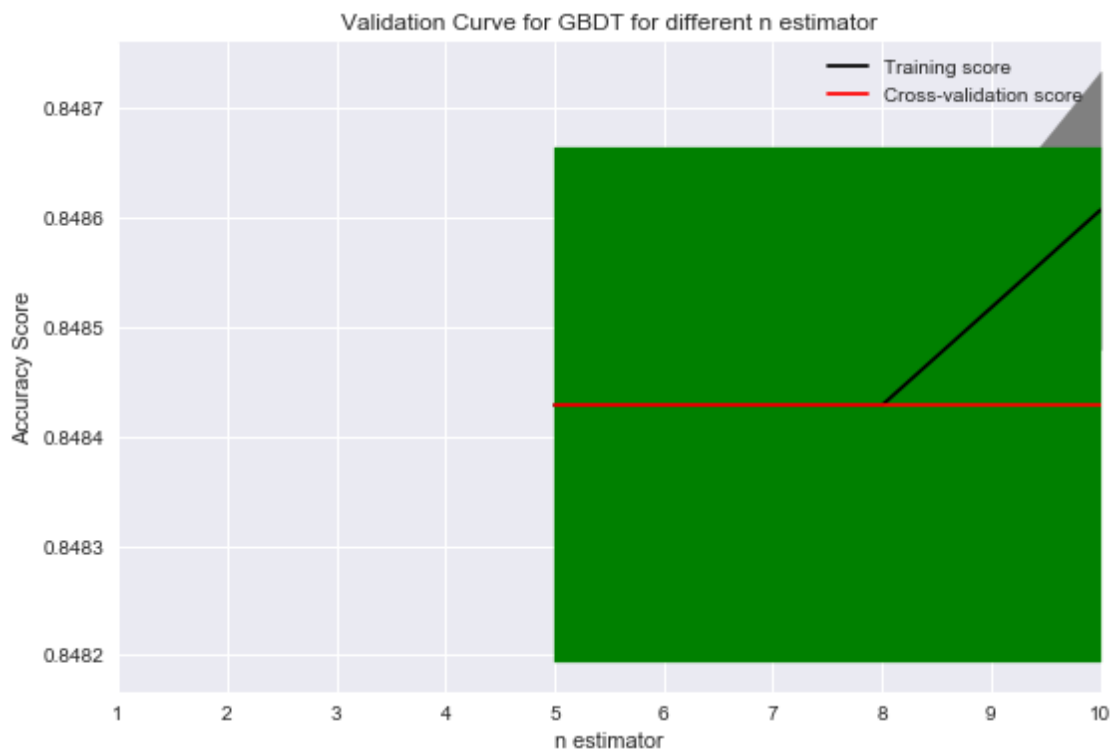
```

param_range=n_estimators,cv=5)

#print(train_scores,test_scores)
train_scores_mean = np.mean(train_scores, axis=1)
train_scores_std = np.std(train_scores, axis=1)
test_scores_mean = np.mean(test_scores, axis=1)
test_scores_std = np.std(test_scores, axis=1)

plt.plot(param_range, train_scores_mean, label="Training score", color="black")
plt.plot(param_range, test_scores_mean, label="Cross-validation score", color="red")
plt.fill_between(param_range, train_scores_mean - train_scores_std, train_scores_mean + train_scores_std, color="black")
plt.fill_between(param_range, test_scores_mean - test_scores_std, test_scores_mean + test_scores_std, color="red")
plt.title("Validation Curve for GBDT for different n estimator")
plt.xlabel("n estimator")
plt.ylabel("Accuracy Score")
plt.xlim(1,10)
plt.tight_layout()
plt.legend(loc="best")
plt.show()

```



```

In [27]: import warnings
warnings.filterwarnings('ignore')
from sklearn.model_selection import validation_curve
#create plot for training and test validation

```

```

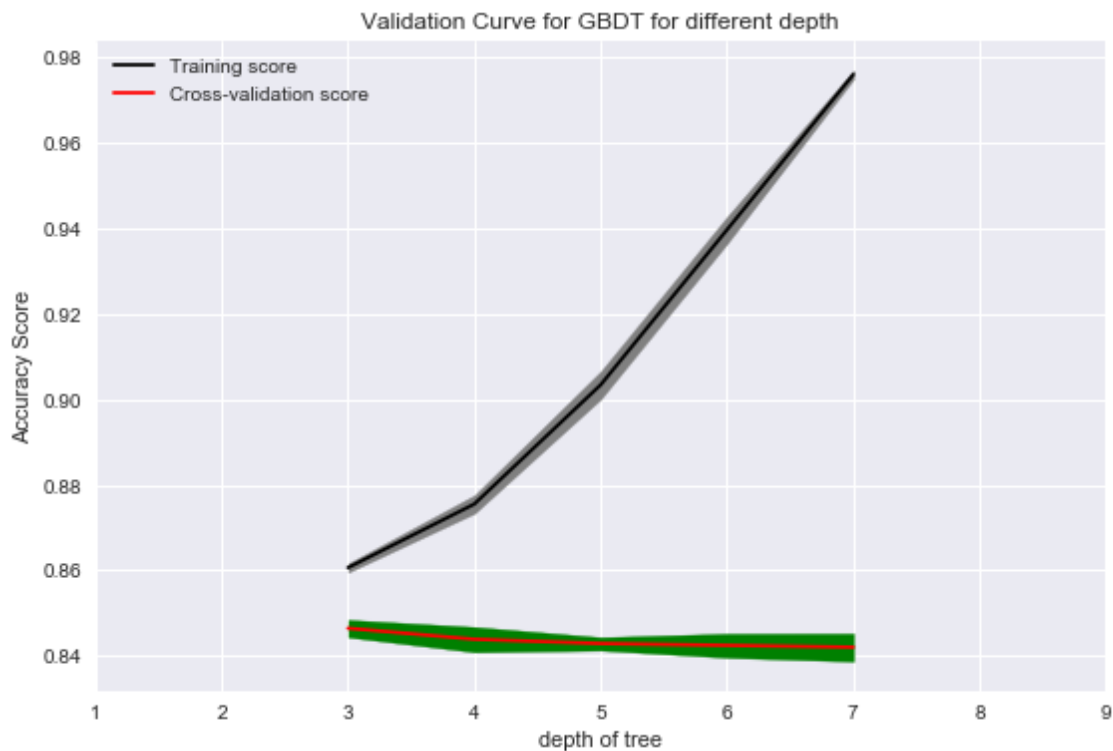
# We cannot put 0 in parameter it will give error while validation_curve
learning_rate=[.9,.8,.1,.2]
n_estimators=[5,8,10]
max_depth=[3,4,5,6,7]
param_range=[3,4,5,6,7]

train_scores, test_scores = validation_curve(GradientBoostingClassifier(), X_train, y,
                                             param_range=max_depth,cv=5)

#print(train_scores,test_scores)
train_scores_mean = np.mean(train_scores, axis=1)
train_scores_std = np.std(train_scores, axis=1)
test_scores_mean = np.mean(test_scores, axis=1)
test_scores_std = np.std(test_scores, axis=1)

plt.plot(param_range, train_scores_mean, label="Training score", color="black")
plt.plot(param_range, test_scores_mean, label="Cross-validation score", color="red")
plt.fill_between(param_range, train_scores_mean - train_scores_std, train_scores_mean + train_scores_std)
plt.fill_between(param_range, test_scores_mean - test_scores_std, test_scores_mean + test_scores_std)
plt.title("Validation Curve for GBDT for different depth")
plt.xlabel("depth of tree")
plt.ylabel("Accuracy Score")
plt.xlim(1,9)
plt.tight_layout()
plt.legend(loc="best")
plt.show()

```



11 Conclusion

GBDT is taking a lot of time for estimator 20/30

12 Steps followed

Only !=3 reviews are taken Mark >3 as positive and <3 as negative. Sort data as per product id in ascending order Deduplication of entries for same profilename,userid, time, text and take first element Get stratified sampling of 10k data Clean html and punctuation Convert to uppercase and word<3 are rejected data sorted on time Split the data in train and test to 70:30

BOW BOW BOW vec created using train data test data is converted using above on train and same applied to test do grid search for different value of learning_rate,n_estimators,max_depth. best model is established with best hyperparameter. model metric is stored in dataframe and crosstable is printed.

TFIDF TFIDF form tfidf vec using train same is used in test to convert rest are same

AVG W2V gensim is used to convert train and test text to W2V AVG seperately rest are same

AVG TFIDF form tfidf vec using train same is used in test to convert gensim and above tfidf is used to convert train and test text to W2V AVG seperately

plot cv error with learning_rate,n_estimators,max_depth

Below is metric

In [29]: aa

```
Out[29]:
```

	accuracy_test	accuracy_train	depth	estimator	fscore_test	fscore_train	\
0	0.864667	0.933714	7	30	0.923367	0.962380	
0	0.866667	0.933714	7	30	0.924699	0.962374	
0	0.828667	0.856143	4	12	0.906033	0.921811	
0	0.831000	0.848714	3	5	0.624343	0.918142	
	learningrate			type			
0	0.2			BOW GDDT			
0	0.2			TFIDF GDDT			
0	0.2			AVG W2V GDDT			
0	0.2			AVG W2V TFIDF GBDT			